


Understanding ARM Instruction Set Architecture (ISA) using ARM Simulator


Dr. Hari T.S. Narayanan



Pre-Requisite




- High School Computer Science
 - Arithmetic and Logic Unit (ALU)
 - Memory, Registers
 - C Programming
 - Compiler
 - Assembly Level Language
 - Assembler
 - Machine Code



Importance of Instruction Set

- Computer Organization course presents the design and development of Computing components (Memory, ALU, I/O Ports, Bus, etc)
- Computer Architecture course on the other hand
 - Defines the interfaces between various components within a Processor
 - Defines programmers' interface to a Processor using Instruction Set Architecture (ISA)



The ARM Instruction Set Architecture

Most of this material is taken from ARM University

Why ARM ISA in Computer Organization Course; that too at the start?

Advanced RISC Machine (ARM)




- ARM family includes 8/16/32/64 bits processors
- Different versions of ARM processors share the same machine Instruction Set
- We will be studying 32-bit ARM processor using a ARM Simulator application, **ARMSim#**
- ARMSim# © R. N. Horspool, W. D. Lyons, M. Serra Department of Computer Science, University of Victoria – Why ARMSIM#?




Main features of the ARM Instruction Set

- All instructions are 32 bits long.
- ARM processor supports 8/16/32/64 bits data values
- ARM processor executes instructions in a 3-Stage pipeline.
- Most of the ARM instructions executed in a single cycle.
- Every instruction can be **conditionally executed**.
- ARM uses load/store architecture (Register-Register Architecture)



Interrupt, Interrupt Handling, and Polling


- In program in execution can be **Interrupted** by an I/O device or a Timer
- CPU suspends the current program*, handles the Interrupting I/O
- Once the I/O is completed CPU continues with the suspended program
- The interrupt handling itself can be Interrupted by another higher priority I/O device
- **Interrupt Handler** is the piece of code that process the I/O or timer
- Interrupt is an efficient alternative to **Polling** or **Busy waiting**
- Interrupt is not the only thing that can alter the flow of execution of a program



Processor Modes

- The ARM has six operating modes:
 - User (unprivileged mode under which most tasks run)
 - FIQ (entered when a high priority (**fast**) **interrupt** occurs)
 - IRQ (entered when a low priority (**normal**) **interrupt** occurs)
 - Supervisor (entered on reset and when a **Software Interrupt instruction** is executed)
 - **Abort** (used to handle memory access violations)
 - **Undef** (used to handle undefined instructions)

Addressing




- ARM includes capability to address byte, half word, and full word (32 bits)
- ARM **instructions are always 32 bits in size**, data can be of 8/16/32 bits
- This means, addresses of instructions start from **0 and increase in steps of 4**
- Instructions are loaded from **Word** boundary – **Instruction Offset**
- **Default Instruction offset is 0**
- **Program Counter** is initialized to 0 when the processor is reset
- The **least significant 2 -bits** of Program Counter are always 0
- This reduces the number of bits required to store **address and address offset of instructions**

Registers



- **Registers** are high speed and low capacity (8/16/32, ...) memory
- ALU operates on data in Data Registers –r0, r1,..., r12
- There are other special purpose Registers; each with specific role
- Examples: **Program Counter (PC)**; **Process Status Register (PSR)**
- Register File and Register banks are physical and logical grouping of Registers respectively
- ARM supports 37 Registers in all

Register Banks



37 Registers

1. 18 General Purpose Registers: r0 – r12, r8fm-r12fm (fm stands of FIQ Mode)
2. 1 Program Counter: pc
3. 6 Stack Pointer (sp): sp-xx
4. 6 Link Register (lr): lr-xx
5. 1 Current Program Status Register: cpsr
6. 5 Saved Program Status Register: spsr-xx

r0-r12, sp-um, lr-um, pc,
cpsr

User mode (um)
Bank – 17 regs

r0-r7, r8fm-r12fm, sp-fm,
lr-fm, pc, cpsr, spsr-fq

FIQ mode (fm)
Bank – 18 regs

r0-r12, sp-im, lr-im, pc,
cpsr, spsr-im

IRQ mode (im)
Bank – 18 regs

r0-r12, sp-sm, lr-sm, pc,
cpsr, spsr-sm

Supervisor mode (sm)
Bank – 18 regs


r0-r12, sp-nm, lr-nm, pc,
cpsr, spsr-nm

Undefined mode (nm)
Bank – 18 regs


r0-r12, sp-am, lr-am, pc,
cpsr, spsr-am

Abort mode (am)
Bank – 18 regs

6 Register Banks



Program Status Registers (PSR)



- Condition code flags

- N = **N**egative result from ALU
- Z = **Z**ero result from ALU
- C = ALU operation **C**arried out
- V = ALU operation o**V**erflowed

Carry versus overflow!

- Interrupt Disable bits.


- I = I: Disables the IRQ.
- F = I: Disables the FIQ.

- Mode bits

- Specify the processor mode

Mode Encoding	
Value	Mode
10000	User
10001	FIQ
10010	IRQ
10011	SVC
10100	Abort
10101	Undef
11111	System

Hands-on Exercises




- Copy ARM Simulator (**Armsim#**) folder
- Install ARM Simulator
- Check if the Simulator installed correctly
- Start the simulator
- Load a sample program (test.s)
- Explore the features of the Simulator



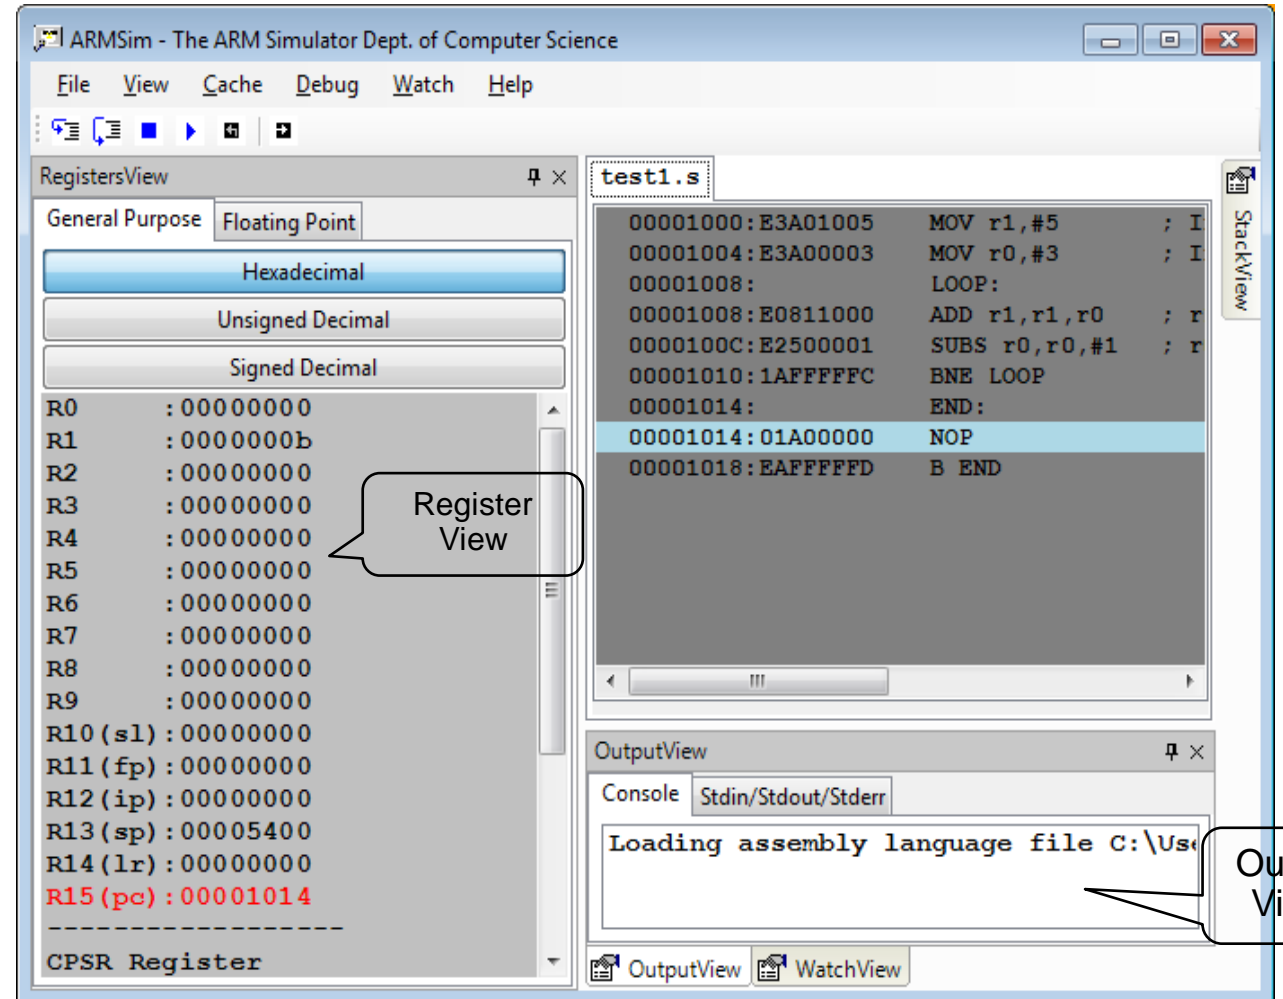
Ex 1: Installing ARMSim


- Install ARMSim 1.9I from ARMSim folder that you copied from my public folder
- Note: This simulator (Windows, Linux and Mac versions available) can also be downloaded from: <http://armsim.cs.uvic.ca/Downloads.html>.
- Override your Anti-virus, if the download fails for some reason
- The file is about 1300 KB
- Follow the default install options. ARMSim installs in few seconds.
- The Simulator is installed in “C:/Program Files (x86)/University of Victoria”
- There will be short cut installed on your desktop
- Launch the simulator to validate your installation



Ex 2: Understanding ARM Simulator UI


- Once you successfully launch ARM Simulator
- Go to “File” menu and using “Load” option, load testI.s file
- Note: If the test.s is not loaded it will not be possible to view the register contents.
- Views
- Registers
- Instructions & Instruction Encoding
- Menu items (Note: Watch & Cache later)





Ex 3: User Mode Bank Registers


- Observe the contents of r0 (R0) to r12 (R12)
- R13 is Stack Pointer (SP)
- R14 is Link Register (LR)
- R15 is Program Counter (PC)
- Current Process Status Register (CPSR) is shown with its bit map
- Correlate R15 (PC) content with Code View instruction address, and memory content in memory view



Exception Handling

- Exceptions are events that interrupt the flow of currently running program
- Exceptions are handled by exception specific code
- This code is loaded (context Switched) and given control after completing the current instruction.
- The **Context Switching** is realized by h/w (processor core)
- Vector table is an important data structure used in Context Switching
- Some Exceptions are: I/O and Timer Interrupts, Reset, SWI

Exception Handling and the Vector Table



- When an exception occurs, the core:

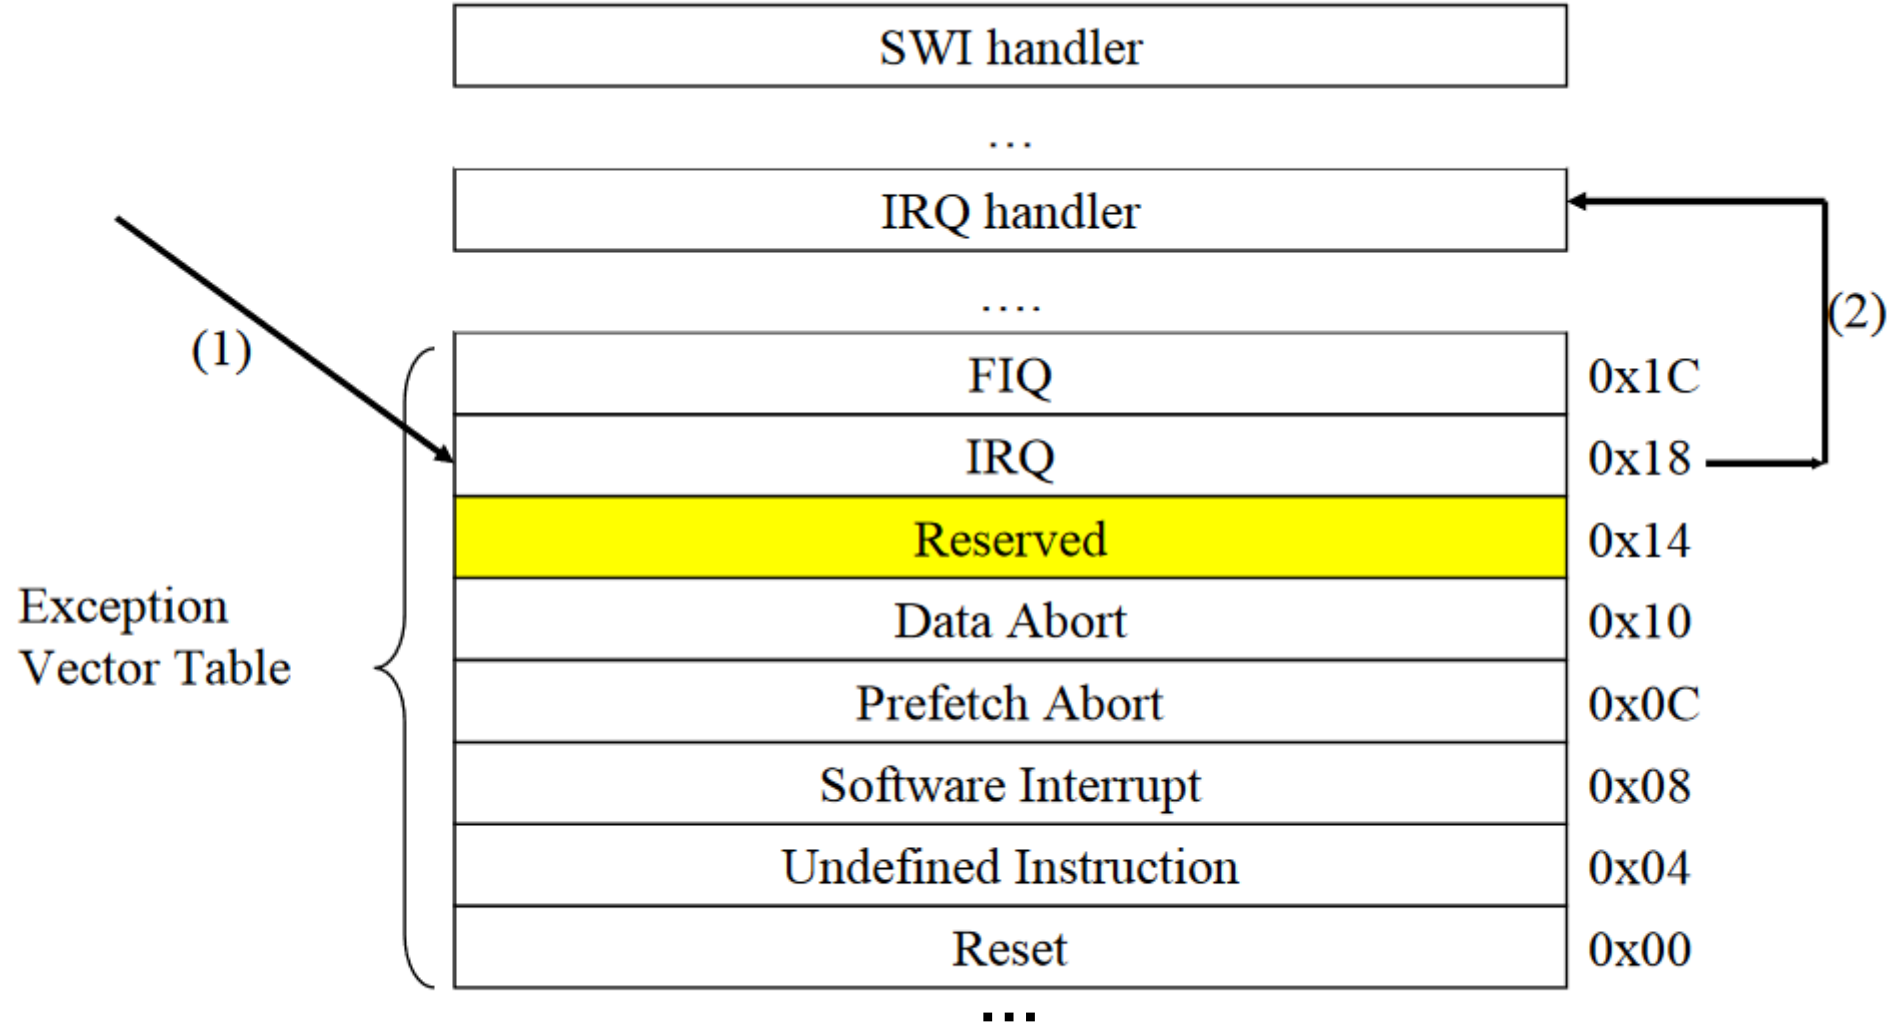

- Copies CPSR into SPSR_<mode>
- Sets appropriate CPSR bits (**especially the mode**)
- Maps in appropriate banked registers
- Stores the “*return address*” in LR_<mode>
- Sets PC to vector address
 - For instance for Software Interrupt PC is set to 0x00000008


- To return, exception handler (sw) needs to:

- Restore CPSR from SPSR_<mode>
- Restore PC from LR_<mode>

0x00000000	Reset	LDR	PC, Reset_Addr
0x00000004	Undefined Instruction	LDR	PC, Undefined_Addr
0x00000008	Software Interrupt	LDR	PC, SWI_Addr
0x0000000C	Prefetch Abort	LDR	PC, Prefetch_Addr
0x00000010	Data Abort	LDR	PC, Abort_Addr
0x00000014	Reserved	NOP	
0x00000018	IRQ	LDR	PC, IRQ_Addr
0x0000001C	FIQ	LDR	PC, FIQ_Addr


Memory Map





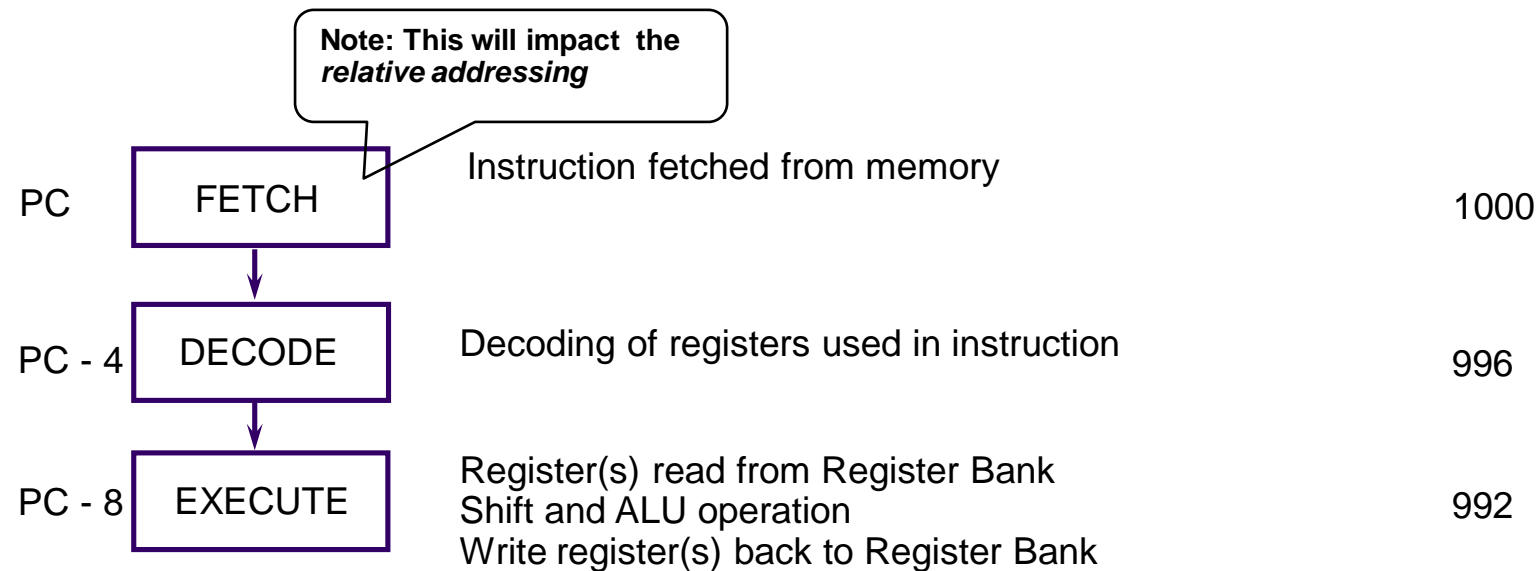
Branching to Absolute Address

- Load the PC with the absolute address of the instruction to be executed!
- This could be done only in certain **modes** – example SWI (Supervisory)
- Example:
 - `ldr pc, = 0X3000`



The Instruction Pipeline

- In order to increase the *instruction throughput* ARM uses a 3-stage pipeline



- PC points to the instruction being fetched, rather than pointing to the instruction being executed.

Branching Instruction with Offset

- The ARM supports B(ranch) instructions with **offset address** (rather than **absolute address**)
- The **target address** to branch is computed at execution using an **offset** (actual offset*4) & **Program Counter (PC)** Value:

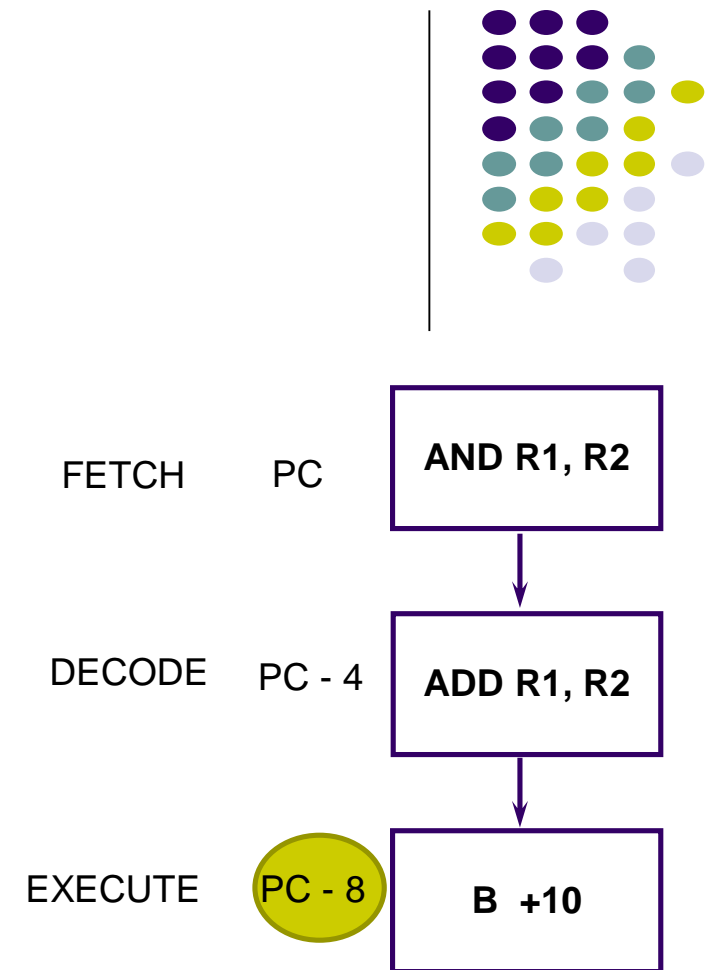
$$\text{target address} = \text{PC} + \text{offset} \times 4;$$

$$= (\text{branch instruction address} + 8) + (\text{offset} \ll 2)$$


- Similarly the *offset address* can be computed from *target* and branch instruction address as follows:

$$\text{Offset} = (\text{target address} - \text{PC}) / 2$$

$$= (\text{target address} - (\text{branch instruction address} + 8)) \gg 2$$



Instruction Types




CPSR

N	Z	C	V																	I	F	T	Mode		
				31	28												8				4				0






Code	Suffix	Flags	Meaning
0000	EQ	Z set	equal
0001	NE	Z clear	not equal
0010	CS	C set	unsigned higher or same
0011	CC	C clear	unsigned lower
0100	MI	N set	negative
0101	PL	N clear	positive or zero
0110	VS	V set	overflow
0111	VC	V clear	no overflow
1000	HI	C set and Z clear	unsigned higher
1001	LS	C clear or Z set	unsigned lower or same
1010	GE	N equals V	greater or equal
1011	LT	N not equal to V	less than
1100	GT	Z clear AND (N equals V)	greater than
1101	LE	Z set OR (N not equal to V)	less than or equal
1110	AL	(ignored)	always



Conditional Execution

- Most processors allow only branches to be executed conditionally.
- However by reusing the condition evaluation hardware, ARM effectively increases number of instructions.
 - All instructions contain a condition field which determines whether the CPU will execute them.
 - Non-executed instructions soak up 1 cycle.
 - Still have to complete cycle so as to allow fetching and decoding of following instructions.
- This removes the need for many branches, which stall the pipeline (3 cycles to refill).
 - Allows very dense in-line code, without branches.
 - The Time penalty of not executing several conditional instructions is frequently less than overhead of the branch or subroutine call that would otherwise be needed.



Using and updating the Condition Field

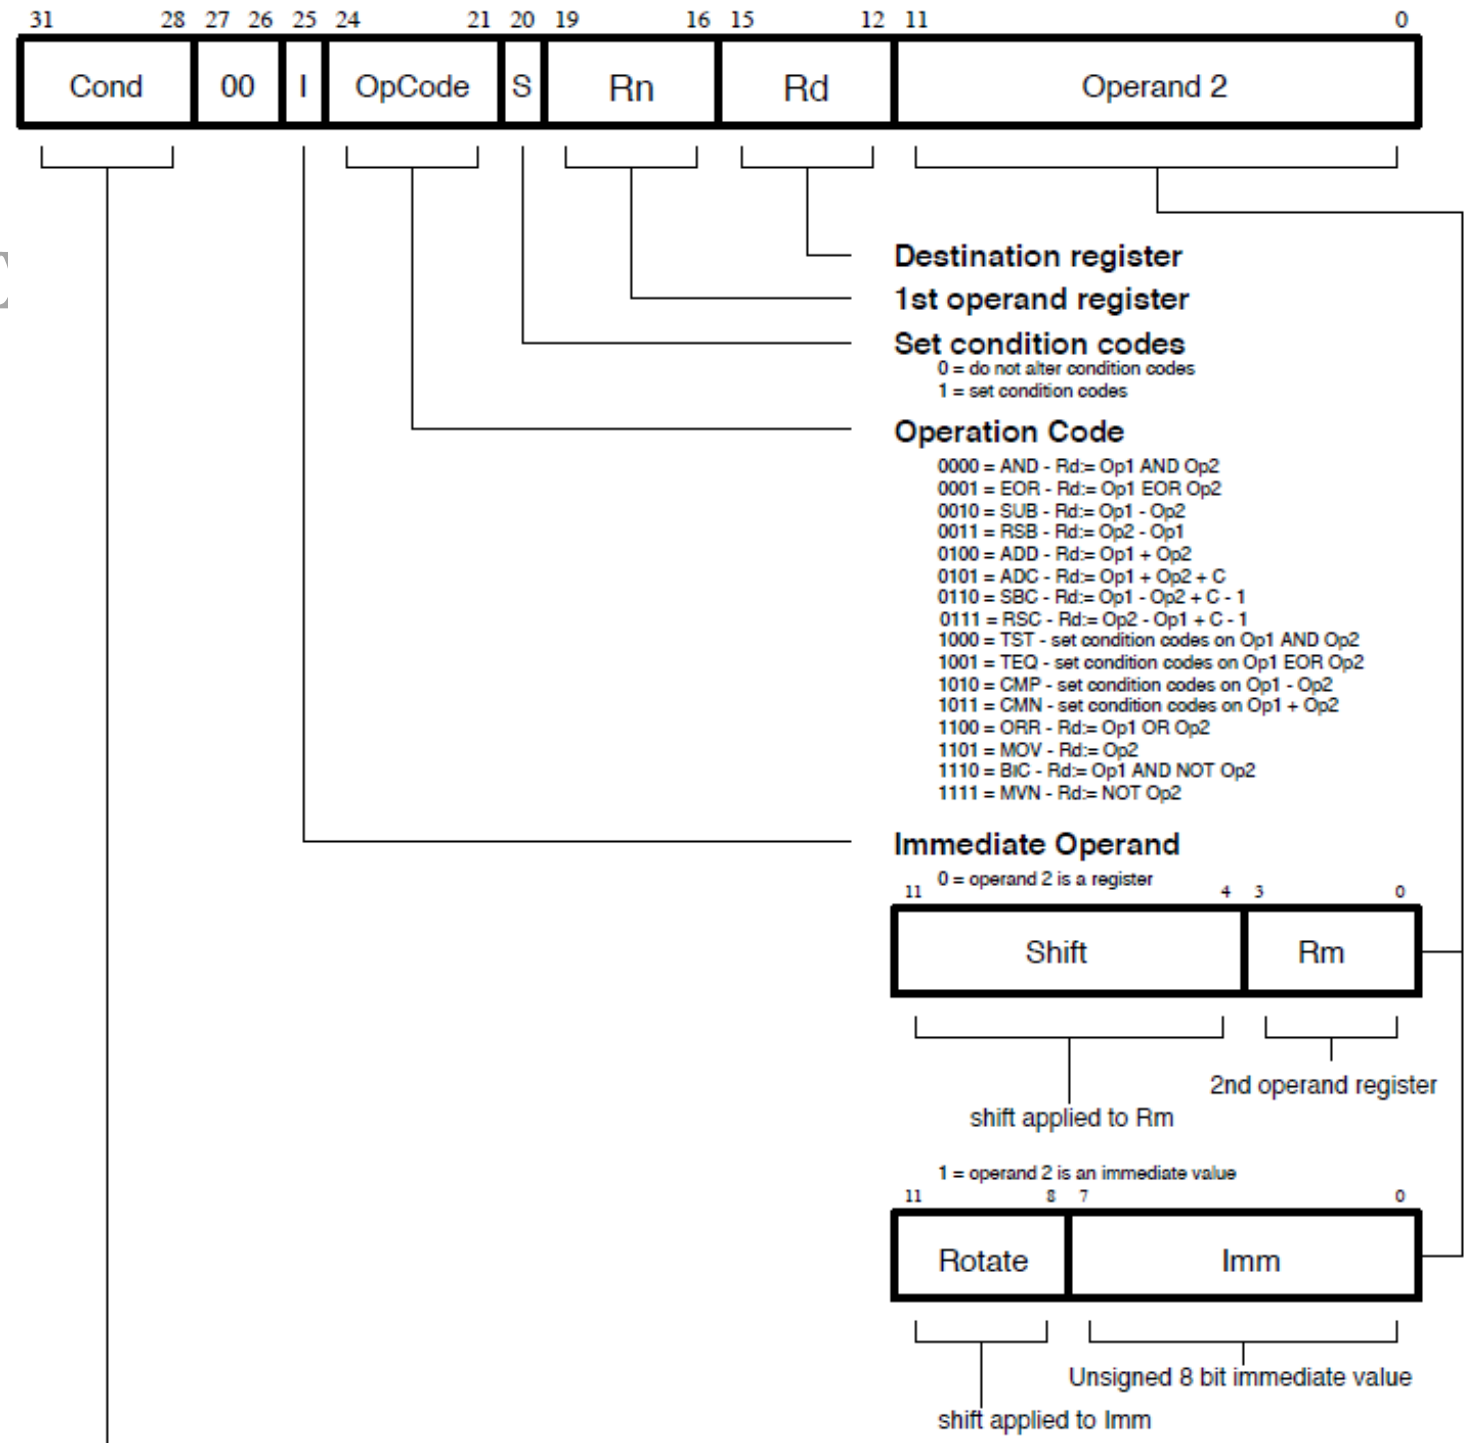
- To execute an instruction conditionally, simply postfix it with the appropriate condition:
 - For example an add instruction takes the form:
 - `ADD r0, r1, r2 ; r0 = r1 + r2 (ADDAL)`
 - To execute this only if the zero flag is set:
 - `ADDEQ r0, r1, r2 ; If zero flag set then r0 = r1 + r2`
- By default, data processing operations do not affect the condition flags (apart from the comparisons where this is the only effect).
- To cause the condition flags to be updated, the S bit of the instruction needs to be set by postfixing the instruction (and any condition code) with an “S”.
 - For example to add two numbers and set the condition flags:
 - `ADDS r0, r1, r2 ;`
 - `r0 = r1 + r2 ; ... and set flags`

✓ Ex 4: PC, Address, Memory


- Add rI, rI, r0

1110 00 0 0100 0 0001 0001 00000000 0000

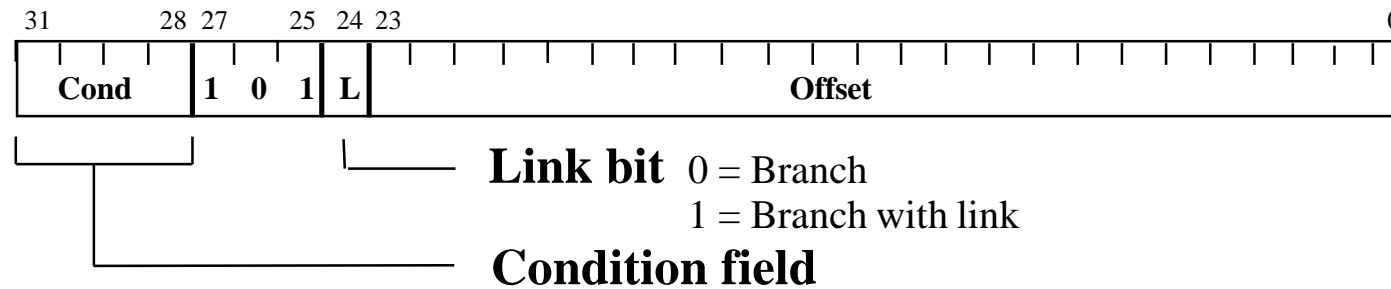
- Binary representation of the above Add instruction




Branch instruction




- `B{<cond>} label ;` versus Branch with Link



- The offset for branch instructions is calculated by the assembler:
 - By taking the difference between the target instruction and the branch instruction address minus 8 (to allow for the pipeline). Then right shift the result by 2 positions.
 - This gives a range of ± 32 Mbytes.





Ex 6: Branch Example (Forward Branch) ✓

- The *target address* to branch is computed as a function *offset address* with respect to *Program Counter (PC) Value*

target address

$$= PC + offset * 4$$

$$= (branch\ instruction\ address + 8) + (offset \ll 2)$$

- Similarly the *offset address* can be computed from *target* and *branch instruction addresses* as follows:

offset

$$= (target\ address - PC) \gg 2$$

$$= (target\ address - (branch\ instruction\ address + 8)) \gg 2$$

- Branch Instruction address: 0x1010; Target Address: 0x1024

Offset = ?; From offset compute Target Address


Now using this offset and Branch instruction address find target address

- Note: use test2.s

test2.s

00001000:E3A01005	MOV r1,#5
00001004:E3A00003	MOV r0,#3
00001008:E3A02008	MOV r2,#8
0000100C:	LOOP:
0000100C:E0821000	ADD r1,r2,r0
00001010:E2500001	SUBS r0,r0,#1
00001014:0AD00003	BEQ END
00001018:EAffFFFFB	B LOOP
0000101C:01A00000	NOP
00001020:01A00000	NOP
00001024:01A00000	NOP
00001028:	END:
00001028:01A00000	NOP
0000102C:01A00000	NOP
00001030:01A00000	NOP
00001034:EAffFFFFB	B END

Test2.s



Ex 7: Branch Example (Backward Branch) ✓

- The *target address* to branch is computed as a function *offset address* with respect to *Program Counter (PC) Value*

target address

$$= PC + (\text{offset address})$$

$$= (\text{branch instruction address} + 8) + (\text{offset} \ll 2)$$

- Similarly the *offset address* can be computed from *target* and *branch instruction addresses* as follows:

offset

$$= (\text{target address} - PC) \ll 2$$

$$= (\text{target address} - (\text{branch instruction address} + 8)) \gg 2$$


- Branch Instruction address: 0x1014; Target Address: 0x1008

Offset = ?; From offset compute Target Address

Now using this offset and Branch instruction address find target address



test2.s


00001000:E3A01005	MOV r1,#5
00001004:E3A00003	MOV r0,#3
00001008:E3A02008	MOV r2,#8
0000100C:	LOOP:
0000100C:E0821000	ADD r1,r2,r0
00001010:E2500001	SUBS r0,r0,#1
00001014:0A000003	BEQ END
00001018:EAF0FFFB	B LOOP
0000101C:01A00000	NOP
00001020:01A00000	NOP
00001024:01A00000	NOP
00001028:	END:
00001028:01A00000	NOP
0000102C:01A00000	NOP
00001030:01A00000	NOP
00001034:EAF0FFFB	B END



Branch instruction with Link

- Branch with Link : `BL{<cond>} sub_routine_label`






Data processing Instructions


- Largest family of ARM instructions, all sharing the same instruction format.
- Contains:
 - Arithmetic operations
 - Comparisons (no results - just set condition codes)
 - Logical operations
 - Data movement between registers
- Remember, this is a load / store architecture
 - These instructions only work on registers, ***NOT*** memory.
- They each perform a specific operation on one or two operands.
 - First operand always a register - Rn
 - Second operand sent to the ALU via barrel shifter.
- We will examine the barrel shifter shortly.

Arithmetic Operations




- Operations are:
 - ADD operand1 + operand2 ; operand1 is Rn
 - ADC operand1 + operand2 + carry
 - SUB operand1 - operand2
 - SBC operand1 - operand2 + carry - 1
 - RSB operand2 - operand1
 - RSC operand2 - operand1 + carry - 1
- Syntax:
 - <Operation>{<cond>}{S} Rd, Rn, Operand2 ; Rd is destination register
- Examples
 - ADD r0, r1, r2
 - SUBGT r3, r3, #1
 - RSBLESS r4, r5, #5

Comparisons




- The only effect of the comparisons is to
 - UPDATE THE CONDITION FLAGS. Thus no need to set S bit.
- Operations are:
 - CMP operand1 - operand2, but result not written
 - CMN operand1 + operand2, but result not written
 - TST operand1 AND operand2, but result not written; **Bitwise AND**
 - TEQ operand1 EOR operand2, but result not written; **Bitwise EXOR**
- Syntax:
 - <Operation>{<cond>} Rn, Operand2
- Examples:
 - CMP r0, r1
 - TSTEQ r2, #5

Logical Operations



- Operations are:
 - AND operand1 AND operand2
 - EOR operand1 EOR operand2
 - ORR operand1 OR operand2
 - BIC operand1 AND NOT operand2 [Bit Clear when operand2 = operand1]
- Syntax:
 - $\langle \text{Operation} \rangle \{ \langle \text{cond} \rangle \} \{ S \} R_d, R_n, \text{Operand2} ; R_d \leftarrow R_n \text{ op Operand2}$
- Examples:
 - AND r0, r1, r2
 - BICEQ r2, r3, #7
 - EORS r1, r3, r0




Data Movement

- Operations are:
 - MOV Rd, operand2
 - MVN Rd, NOT operand2


Note that these instructions make no use of operand1.

- Syntax:
 - <Operation> {<cond>} {S} Rd, Operand2
- Examples:
 - MOV r0, r1
 - MOVS r2, #10
 - MVNEQ r1, #0

The Barrel Shifter




- The ARM doesn't have actual shift instructions.
- Instead it has a barrel shifter which provides a mechanism to carry out shifts as part of other **instructions with second operand**
- Barrel shifter performs shift operations faster than regular shifter
- So what operations does the barrel shifter support?



Barrel Shifter - Left Shift

- Shifts left by the specified amount (multiplies by powers of two) e.g.
LSL #5 = multiply by 32
- CF is Carry Flag bit **Logical Shift Left (LSL)**





Barrel Shifter - Right Shifts

Shift Right - Logical

Shifts right by the specified unsigned amount
(divides by powers of two) e.g.

LSR #2 = divide by 4


	7	6	5	4	3	2	1	0	Comment
	0	0	1	0	0	0	0	0	Representing 32
									Do LSR 1
0 →	0	0	0	1	0	0	0	0	This is 16
									Do LSR 1
0 →	0	0	0	0	1	0	0	0	This is 8

Shift Right - Arithmetic

Shifts right (divides by powers of two) and
preserves the sign bit, for 2's complement
operations. e.g.

ASR #2 = divide by 4


	7	6	5	4	3	2	1	0	Comment
	1	1	1	0	0	0	0	0	Representing -32 in 2s complement
									Do ASR 1
	1	1	1	1	0	0	0	0	This is -16 in 2s complement form
									Do ASR 1
	1	1	1	1	1	0	0	0	This is -8 in 2s complement form



Barrel Shifter - Rotations


Rotate Right (ROR)


- Similar to an ASR but the bits wrap around as they leave the LSB and appear as the MSB.
e.g. ROR #4
- Note the last bit rotated is also used as the Carry Out.



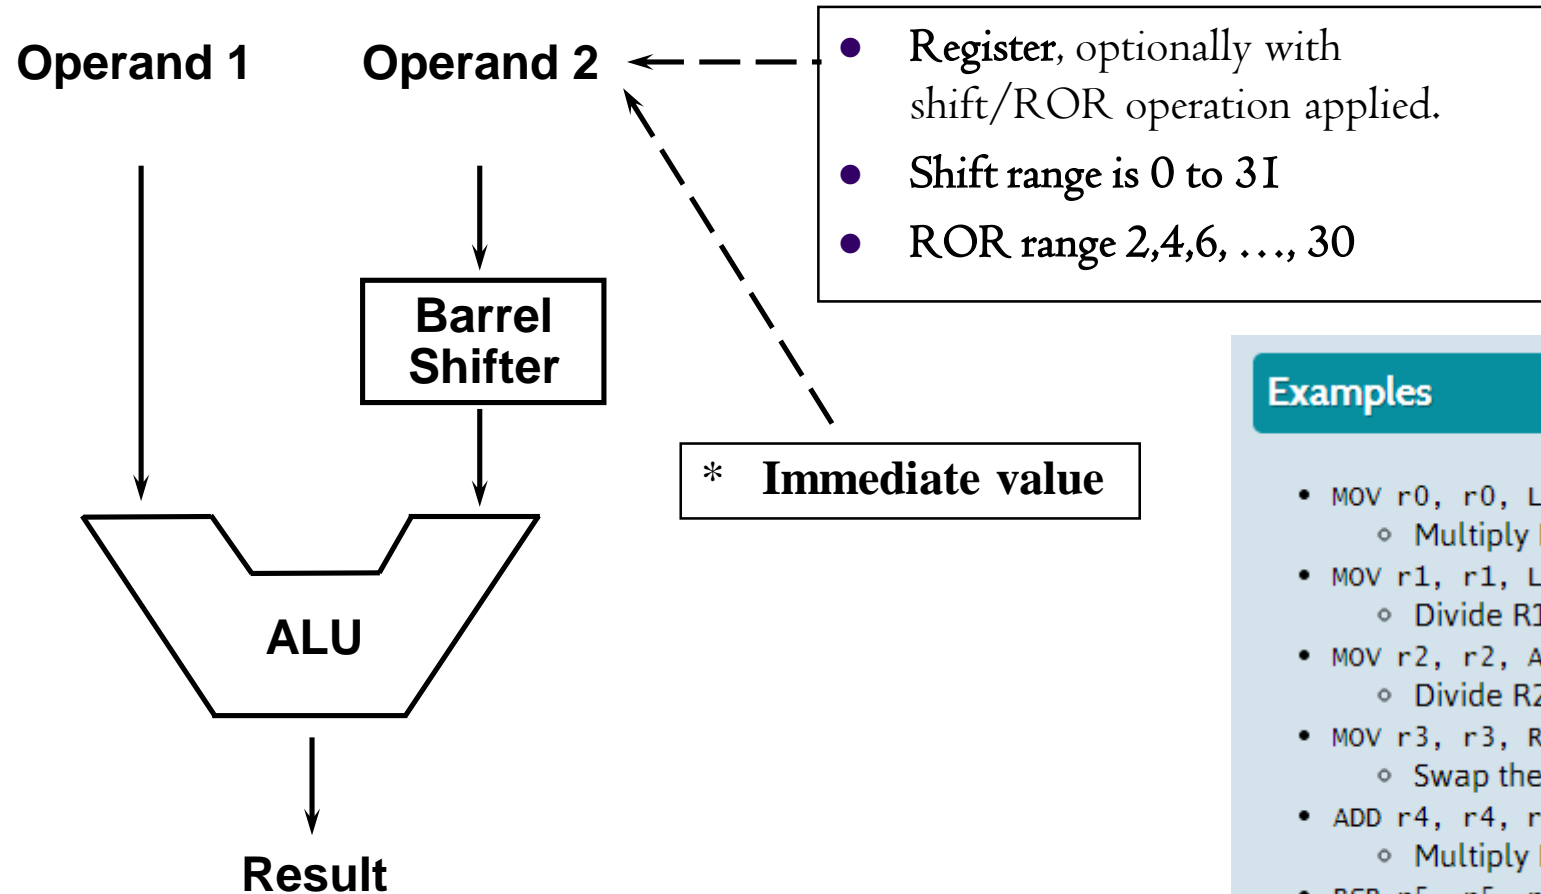
Rotate Right Extended (RRX)

- This operation uses the CPSR C flag as a 33rd bit.
- Rotates right by 1 bit. Encoded as ROR #0.






Using the Barrel Shifter: The Second Operand







Examples

- `MOV r0, r0, LSL #1`
 - Multiply R0 by two.
- `MOV r1, r1, LSR #2`
 - Divide R1 by four (unsigned).
- `MOV r2, r2, ASR #2`
 - Divide R2 by four (signed).
- `MOV r3, r3, ROR #16`
 - Swap the top and bottom halves of R3.
- `ADD r4, r4, r4, LSL #4`
 - Multiply R4 by 17. ($N = N + N * 16$)
- `RSB r5, r5, r5, LSL #5`
 - Multiply R5 by 31. ($N = N * 32 - N$)

Ex - Using Shifted Register ✓



- Using a multiplication instruction to multiply by a constant means first loading the constant into a register and then waiting a number of internal cycles for the instruction to complete.
- A more optimum solution can often be found by using some combination of MOVs, ADDs, SUBs and RSBs with shifts.
 - Multiplications by a constant equal to a $((\text{power of } 2) \pm 1)$ can be done in one cycle.
- Example: $r0 = r1 * 5$; **MUL R1, R1, #5**
 $= r1 + (r1 * 4)$
 **ADD r0, r1, r1, LSL #2**
- Example: $r2 = r3 * 105$
 $= r3 * 15 * 7$
 $= r3 * (16 - 1) * (8 - 1)$
 **RSB r2, r3, r3, LSL #4** ; $r2 = r3 * 15$
 **RSB r2, r2, r2, LSL #3** ; $r2 = r2 * 7$




Large constants with limited size of Immediate Value

- Immediate operand size is 12 bits.
- Out of this, the least significant 8 bits are used for constant value
- Rest of the 4 bits are used of ROR
- Value 1 in these bits code ROL 2
- Value 2 in these bits code ROL 4
- ...
- Value 15 in these bits code ROL 30

Second Operand: Immediate Value (2)


- **This gives us:**
 - 0 - 255 [0 - 0xff]
 - 256,260,264,...,1020 [0x100-0x3fc, step 4, 0x40-0xff ror 30]
 - 1024,1040,1056,...,4080 [0x400-0xff0, step 16, 0x40-0xff ror 28]
 - 4096,4160, 4224,...,16320 [0x1000-0x3fc0, step 64, 0x40-0xff ror 26]
- **These can be loaded using, for example:**
 - `MOV r0, #0x40, 26` ; => `MOV r0, #0x1000` (ie 4096)
- **To make this easier, the assembler will convert to this form for us if simply given the required constant:**
 - `MOV r0, #4096` ; => `MOV r0, #0x1000` (ie 0x40 ror 26)
- **The bitwise complements can also be formed using MVN:**
 - `MOV r0, #0xFFFFFFFF` ; assembles to `MVN r0, #0`
- **If the required constant cannot be generated, an error will be reported.**

Notes ✓




- ROR 30 is same ROL (32-30)
- ROR 0 gives us constants 0 – 255 (0000 0000 to 1111 1111)
- ROR 30 or ROL 2 gives constants 256, 260, ... 1020 (0100 0000 to 0xFF)
- ROR 28 or ROL 4 gives constants 1024, 1040, ... , 4080
- ROR 26 or ROL 6 gives constants 4096, 4128, ... , 10208
- And so on ... till ROR 2

Load / Store Instructions





- The ARM is a Load / Store Architecture:
 - Does not support memory to memory data processing operations.
 - Must move data values into registers before using them.
- This might sound inefficient, but in practice isn't:
 - Load data values from memory into registers.
 - Process data in registers using a number of data processing instructions which are not slowed down by memory access.
 - Store results from registers out to memory.
- The ARM has three sets of instructions which interact with main memory. These are:
 - Single register data transfer (LDR / STR).
 - Block data transfer (LDM/STM).
 - Single Data Swap (SWP).

Load and Store



- The memory location to be accessed is held in a base register
 - STR r0, [r1] ; Store contents of r0 to location pointed to ; by contents of r1.
 - LDR r2, [r1] ; Load r2 with contents of memory location ; pointed to by contents of r1.







Load/Store variations & Syntax

- The basic load and store instructions are:
 - Load and Store Word or Byte
 - LDR / STR / LDRB / STRB
- ARM Architecture Version 4 also adds support for half words and signed data.
 - Load and Store Half word
 - LDRH / STRH
 - Load Signed Byte or Half word - load value and sign extend it to 32 bits.
 - LDRSB / LDRSH
- All of these instructions can be conditionally executed by inserting the appropriate condition code after STR / LDR.
 - e.g. LDREQB
- Syntax:
 - `<LDR | STR> {<cond>} {<size>} Rd, <address>`

Exercise




- Write ARM code snippet that store some data in 5 consecutive half words starting from memory address 0x2000 and retrieve those 5 half words and print the content on stdout



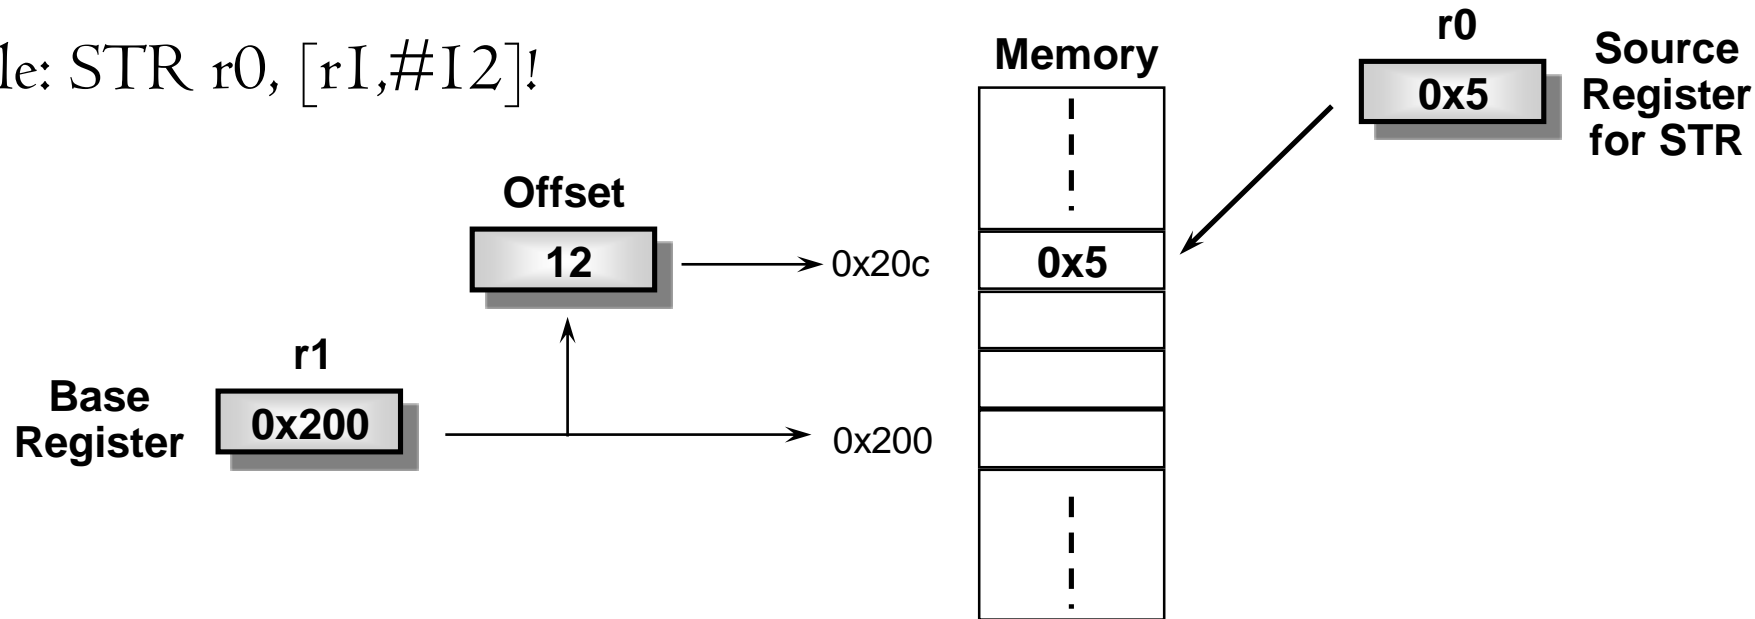
Load and Store with Indexing

- As well as accessing the actual location contained in the base register, these instructions can access a location offset from the base register pointer.
- This **offset** can be
 - An unsigned 12bit immediate value (ie 0 - 4095 bytes).
 - A register, optionally shifted by an immediate value
- This can be either added or subtracted from the base register:
 - Prefix the offset value or register with '+' (default) or '-'.
- This **offset** can be applied:
 - before the transfer is made: *Pre-indexed addressing*
 - optionally *auto-incrementing* the base register, by postfixing the instruction with an '!'.
 - after the transfer is made: *Post-indexed addressing*
 - causing the base register to be *auto-incremented*.




Load and Store Word Pre-indexed Addressing

- Example: STR r0, [r1, #12]!

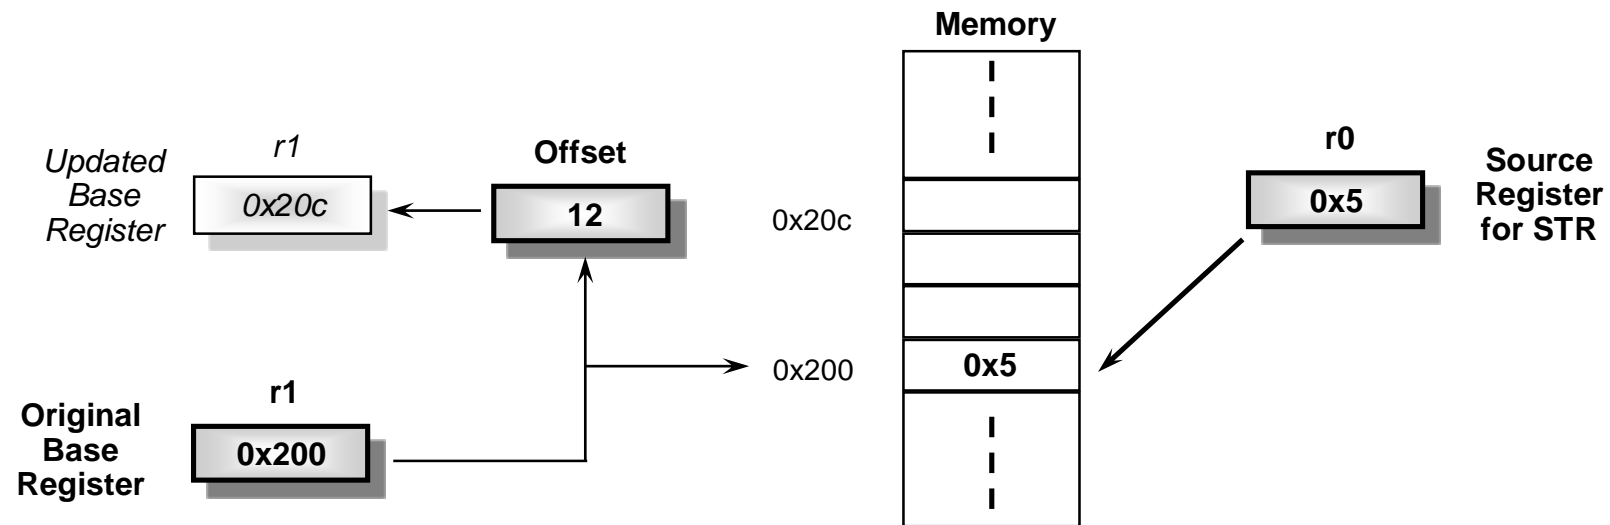


- To auto-pre-increment base pointer to `0x20c` use: `STR r0, [r1, #12]!`
- To auto-pre-decrement base pointer to `0x1f4` use: `STR r0, [r1, #-12]!`
- Immediate operand can be a shifted Register!




Load and Store Word Post-indexed Addressing

- Example: STR r0, [r1], #12




- To auto-post-increment the base register to location `0x1f4` instead use:
 - STR r0, [r1], #-12

Code Snippets Pre-Indexing and Post-Indexing




- Develop ARM assembly code snippet that does Pre-indexing and Post-indexing. Demonstrate this with ARM Simulator
- Use Data Cache for various Block size to show how mis-rate decreases with larger block size
- Also justify the hit and miss count for instruction cache



Swap Instructions



- Atomic operation of a memory read followed by a memory write which moves byte or word quantities between registers and memory.
- Syntax:
 - $\text{SWP}\{\text{<cond>}\}\{\text{B}\}\text{Rd}, \text{Rm}, [\text{Rn}]$



- Thus to implement an actual swap of contents make $\text{Rd} = \text{Rm}$.
- The compiler cannot produce this instruction.


Interrupting an Instruction

Software Interrupt (SWI)




- In effect, a SWI is a user-defined instruction.
- It causes an exception trap to the SWI hardware vector (thus causing a change to supervisor mode, plus the associated state saving), thus causing the SWI exception handler to be called.
- The handler can then examine the comment field of the instruction to decide what operation has been requested.
- By making use of the SWI mechanism, an operating system can implement a set of privileged operations which applications running in user mode can request.
- See Exception Handling Module for further details.

Exception Handling and the Vector Table



- When an exception occurs, the core:
 - Copies CPSR into SPSR_<mode>
 - Sets appropriate CPSR bits
 - Maps in appropriate banked registers
 - Stores the “*return address*” in LR_<mode>
 - Sets PC to vector address
 - For instance for *Undefined instruction* PC is set to 0x00000008
- To return, exception handler needs to:
 - Restore CPSR from SPSR_<mode>
 - Restore PC from LR_<mode>

0x00000000	Reset
0x00000004	Undefined Instruction
0x00000008	Software Interrupt
0x0000000C	Prefetch Abort
0x00000010	Data Abort
0x00000014	Reserved
0x00000018	IRQ
0x0000001C	FIQ




Assembler Directives

- Hello: **.asciz** "Hello Class\n"
- **.equ** SEG_A, 0x80
- **.equ** Ten, 10
- **.set** counter, 1
- **.set** counter, counter+1
- Address: **.word** 0
- **.align**

Reference:

<https://community.arm.com/processors/blog/posts/useful-assembler-directives-and-macros-for-the-gnu-assembler>



SWI – ARMSIM Specific I/O

- There are two plugins with ARMSIM

1. Basic I/O Plugin


- File I/O
- Memory management

2. Embest Board Plugin

1. LED
2. Keypad
3. Buttons
4. LED Number Segment

The I/O with SWI is unique to ARMSim – this is NOT a standard practice

Printing a Character & a string to STDOUT ✓



- **Hello:** .asciz "Hello Class!\n"
- mov r0, #'I'
- swi 0x00
- mov r0, #':'
- swi 0x00
- ldr r0, =Hello
- swi 0x02
- **End:**
- b End




SWI Comment Field Encoding for ARMSim

Opcode	Description and Action	Inputs	Outputs	EQU
swi 0x00	Display Character on Stdout	r0: the character		SWI_PrChr
swi 0x02	Display String on Stdout	r0: address of a null terminated ASCII string	(see also 0x69 below)	
swi 0x11	Halt Execution			SWI_Exit
swi 0x12	Allocate Block of Memory on Heap	r0: block size in bytes	r0:address of block	SWI_MeAlloc
swi 0x13	Deallocate All Heap Blocks			SWI_DAlloc
swi 0x66	Open File (mode values in r1 are: 0 for input, 1 for output, 2 for appending)	r0: file name, i.e. address of a null terminated ASCII string containing the name r1: mode	r0:file handle If the file does not open, a result of -1 is returned	SWI_Open
swi 0x68	Close File	r0: file handle		SWI_Close
swi 0x69	Write String to a File or to Stdout	r0: file handle or Stdout r1: address of a null terminated ASCII string		SWI_PrStr



SWI Comment Field Encoding for ARMSim

Opcode	Description and Action	Inputs	Outputs	EQU
<code>swi 0x6a</code>	Read String from a File	r0: file handle r1: destination address r2: max bytes to store	r0: number of bytes stored	<code>SWI_RdStr</code>
<code>swi 0x6b</code>	Write Integer to a File	r0: file handle r1: integer		<code>SWI_PrInt</code>
<code>swi 0x6c</code>	Read Integer from a File	r0: file handle	r0: the integer	<code>SWI_RdInt</code>
<code>swi 0x6d</code>	Get the current time (ticks)		r0: the number of ticks (milliseconds)	<code>SWI_Timer</code>



File Open & Close

```
InFileName: .asciz  "Infile1.txt"
InFileError: .asciz  "Unable to open input file\n"
            .align
InFileHandle: .word  0
```


The following lines of code open the file called "Infile1.txt" for input and store its file "handle", returned in R0 by the opening call, into the appropriate memory location:

```
ldr r0,=InFileName      @ set Name for input file
mov r1,#0                @ mode is input
swi SWI_Open             @ open file for input
bcs InFileError          @ if error?
ldr r1,=InFileHandle     @ load input file handle
str r0,[r1]              @ save the file handle
```


File Close

```
load the file handle
ldr      r0,=InFileHandle
ldr      r0,[r0]
swi      SWI_Close
```

File Input




```
    ldr  r0,=InFileHandle
    ldr  r0,[r0]
    ldr  r1,=CharArray
    mov  r2,#80
    swi  0x6a
    bcs  ReadError
    ...
InFileHandle: .word 0
CharArray: .skip 80
```



Exercise

- Read Section 9 of ARMSim User Guide with objective to write code for Embest Board Plugin
- This section will not be taught in the class – you are on your own!
- A similar exercise on Cache performance measurement with ARMSIM will be given in four weeks now.



Multiplication Instruction

- ARM implement Booth's algorithm for multiplication
- **Long** and **Accumulate** options are supported
- Syntax for non-long multiplication:

- **Multiply**

- `MUL{<cond>}{S} Rd, Rm, Rs` ; $Rd = Rm * Rs$


- **Multiply Accumulate** - **does addition for free**

- `MLA{<cond>}{S} Rd, Rm, Rs, Rn` ; $Rd = (Rm * Rs) + Rn$

- **Restrictions on use:**

- **Rd and Rm cannot be the same register**


- **Can be avoided by swapping Rm and Rs around. This works because multiplication is commutative.**



Syntax for Long Multiplication ✓

- 64 bit result
 - Unsigned Long and Accumulate
 - Signed Long and Accumulate
- **UMULL{<cond>}{S} RdLo,RdHi,Rm,Rs**
 - **UMLAL{<cond>}{S} RdLo,RdHi,Rm,Rs**
 - **SMULL{<cond>}{S} RdLo, RdHi, Rm, Rs**
 - **SMLAL{<cond>}{S} RdLo, RdHi, Rm, Rs**


Using Link Register



```
start: MOV    r0, #I0        ; Set up parameters
      MOV    rI, #3
      BL     add             ; Call subroutine

ret:   ...
      ...
      ...
      SWI    0xII
add:  ADD    r0, r0, rI      ; Subroutine code
      BX     lr             ; Return from subroutine
```

Using Stack



```
subroutine  PUSH    {r5-r7,lr} ; Push working registers and lr
            ; code
            BL      somewhere_else
            ; code
            POP     {r5-r7,pc} ; Pop working registers and pc
```

What is not discussed?

1. Block Transfer
2. Writing C code, Compiling it for ARM
3. ARM and Thumb coexistence

